

# Server Capacity Analysis & Optimization Report

## Super Mario Quiz Quest

The screenshot shows the 'Preview Server settings' section of the Netlify CLI configuration. It includes fields for 'Preview Server command' (Not set), 'Target port' (Not set), and 'Preview Server size' (1 vCPU, 4 GB RAM, 20 GB Storage). A 'Learn more about configuring the Netlify CLI' link and a 'Configure' button are also visible.

## Executive Summary

Your Super Mario Quiz Quest game is deployed on **Netlify's Preview Server** with the following specifications:

- **1 vCPU**
- **4 GB RAM**
- **20 GB Storage**

**Key Finding:** This is a **client-side rendered (CSR) React application** where 95%+ of game logic runs in the user's browser. The server primarily serves static assets and handles database queries through Supabase.

## 1. Concurrent User Capacity Analysis

### Current Architecture Assessment

Your application follows a **JAMstack architecture**:

- **Frontend:** React SPA built with Vite
- **Backend:** Supabase (PostgreSQL) - separate managed service
- **Hosting:** Netlify (static file CDN)
- **Game Engine:** Client-side Canvas rendering

### Capacity Breakdown

#### Static Asset Serving (Netlify CDN)

- **Estimated Capacity: 1,000-5,000+ concurrent users**

- **Rationale:**
  - Static files (HTML, CSS, JS) are served from Netlify's global CDN
  - Minimal server processing required
  - Your built application is ~2-5 MB (estimated)
  - With 4GB RAM, can easily cache and serve thousands of concurrent requests

## Database Queries (Supabase)

- **Estimated Capacity: 500-2,000 concurrent users**
- **Rationale:**
  - Each game session makes 1-3 database queries:
    1. Initial question load (on game start)
    2. Potential additional queries if difficulty changes
  - Supabase free tier handles ~500 concurrent connections
  - Your queries are simple SELECT statements with indexes
  - No real-time subscriptions or complex joins

## Client-Side Game Processing

- **Estimated Capacity: Unlimited** (runs on user's device)
- **Rationale:**
  - All game logic (physics, collision detection, rendering) runs in browser
  - Zero server load for gameplay
  - Each user's device handles their own game instance

## Overall Concurrent User Estimate

**Conservative Estimate: 500-1,000 concurrent players**

**Optimistic Estimate: 2,000-3,000 concurrent players**

**Bottleneck:** Supabase database connections, not your Netlify server.

---

## 2. Optimization Recommendations

### 🎯 High Impact Optimizations

#### A. Implement Question Caching (Frontend)

**Problem:** Currently fetching questions from Supabase on every game start.

**Solution:** Cache questions in browser storage.

```
// Add to databaseService.ts
const CACHE_KEY = 'quiz_questions_cache';
const CACHE_DURATION = 24 * 60 * 60 * 1000; // 24 hours

interface CachedQuestions {
```

```

data: Question[];
timestamp: number;
topic: string;
difficulty?: string;
}

export async function getRandomQuestionsWithCache(
  count: number,
  topic?: string,
  difficulty?: 'easy' | 'medium' | 'hard'
): Promise<Question[]> {
  const cacheKey = `${CACHE_KEY}_${topic}_${difficulty || 'all'}`;

  // Try to get from localStorage
  const cached = localStorage.getItem(cacheKey);
  if (cached) {
    const parsed: CachedQuestions = JSON.parse(cached);
    const age = Date.now() - parsed.timestamp;

    if (age < CACHE_DURATION && parsed.data.length >= count) {
      // Return random subset from cache
      return parsed.data
        .sort(() => Math.random() - 0.5)
        .slice(0, count);
    }
  }

  // Fetch from database
  const questions = await getRandomQuestions(count * 3, topic, difficulty);

  // Cache the results
  const cacheData: CachedQuestions = {
    data: questions,
    timestamp: Date.now(),
    topic: topic || 'all',
    difficulty
  };
  localStorage.setItem(cacheKey, JSON.stringify(cacheData));

  return questions.slice(0, count);
}

```

## Impact:

- Reduces database queries by **80-90%**
  - Increases capacity to **5,000-10,000 concurrent users**
  - Faster game start times
- 

## B. Prefetch Questions on Topic Selection

**Problem:** Database query happens after user clicks "Start Game", causing delay.

**Solution:** Fetch questions immediately when topic is selected.

```
// In GameUI.tsx
const [prefetchedQuestions, setPrefetchedQuestions] = useState<Question[] | null>(null);

useEffect(() => {
  if (topic && !prefetchedQuestions) {
    // Prefetch questions in background
    getRandomQuestionsWithCache(20, topic, difficulty)
      .then(setPrefetchedQuestions)
      .catch(console.error);
  }
}, [topic, difficulty]);

// Pass prefetchedQuestions to handleStart
```

### Impact:

- Eliminates perceived loading time
  - Better user experience
  - Spreads database load over time
- 

## C. Optimize Database Query Strategy

**Current Issue:** Fetching 3x questions then shuffling client-side is inefficient.

**Solution:** Use PostgreSQL's `RANDOM()` function with proper indexing.

```
// Update getRandomQuestions in databaseService.ts
export async function getRandomQuestions(
  count: number,
  topic?: string,
  difficulty?: 'easy' | 'medium' | 'hard'
): Promise<Question[]> {
  const client = getSupabaseClient();

  let query = client
    .from('questions')
    .select('*')
    .limit(count);

  if (topic) {
    query = query.eq('topic', topic);
  }

  if (difficulty) {
    query = query.eq('difficulty', difficulty);
  }
```

```
// Use Supabase's RPC for random selection
const { data, error } = await client.rpc('get_random_questions', {
  p_count: count,
  p_topic: topic,
  p_difficulty: difficulty
});

if (error) {
  console.error('Error fetching random questions:', error);
  throw new Error(`Failed to fetch random questions: ${error.message}`);
}

return (data as QuestionRow[]).map(rowToQuestion);
}
```

Add to schema.sql:

```
CREATE OR REPLACE FUNCTION get_random_questions(
  p_count INT,
  p_topic VARCHAR DEFAULT NULL,
  p_difficulty VARCHAR DEFAULT NULL
)
RETURNS SETOF questions AS $$ 
BEGIN
  RETURN QUERY
  SELECT * FROM questions
  WHERE (p_topic IS NULL OR topic = p_topic)
    AND (p_difficulty IS NULL OR difficulty = p_difficulty)
  ORDER BY RANDOM()
  LIMIT p_count;
END;
$$ LANGUAGE plpgsql;
```

**Impact:**

- Reduces data transfer by **66%**
  - Faster query execution
  - Lower memory usage
- 

#### D. Implement Service Worker for Asset Caching

**Solution:** Add Progressive Web App (PWA) capabilities.

```
// Create public/sw.js
const CACHE_NAME = 'mario-quiz-v1';
const urlsToCache = [
  '/',
  '/index.html',
  '/css/style.css',
  '/js/script.js',
  '/images/icon.png'
```

```
'/index.html',
'/assets/index.css',
'/assets/index.js'
];

self.addEventListener('install', (event) => {
  event.waitUntil(
    caches.open(CACHE_NAME)
      .then((cache) => cache.addAll(urlsToCache))
  );
});

self.addEventListener('fetch', (event) => {
  event.respondWith(
    caches.match(event.request)
      .then((response) => response || fetch(event.request))
  );
});
```

## Impact:

- Offline gameplay capability
  - Zero server requests for returning users
  - Instant load times
- 

## 🔧 Medium Impact Optimizations

### E. Lazy Load Audio Assets

**Current:** All audio files loaded upfront.

**Solution:** Load audio on-demand.

```
// In audioService.ts
private loadedSounds = new Set<string>();

async play(soundName: string) {
  if (!this.loadedSounds.has(soundName)) {
    await this.loadSound(soundName);
    this.loadedSounds.add(soundName);
  }
  // Play sound
}
```

### F. Optimize Canvas Rendering

**Current:** Rendering entire canvas every frame.

**Solution:** Implement dirty rectangle rendering.

```
// Track changed regions
const dirtyRegions: Rectangle[] = [];

function render() {
  if (dirtyRegions.length === 0) return; // Skip if nothing changed

  dirtyRegions.forEach(region => {
    ctx.clearRect(region.x, region.y, region.width, region.height);
    // Render only this region
  });

  dirtyRegions.length = 0;
}
```

---

## G. Database Connection Pooling

**Current:** Supabase client creates new connections.

**Solution:** Already handled by Supabase, but ensure proper client reuse.

```
// Ensure single Supabase client instance (already implemented)
let supabase: SupabaseClient | null = null;

function getSupabaseClient(): SupabaseClient {
  if (!supabase) {
    supabase = createClient(supabaseUrl, supabaseAnonKey);
  }
  return supabase;
}
```

---

## III Low Impact Optimizations

### H. Code Splitting

```
// vite.config.ts
export default defineConfig({
  build: {
    rollupOptions: {
      output: {
        manualChunks: {
          'react-vendor': ['react', 'react-dom'],
          'supabase-vendor': ['@supabase/supabase-js'],
        }
      }
    }
  }
})
```

```
    }
}
});
```

## I. Image Optimization

- Convert images to WebP format
- Use responsive images with `srcset`
- Lazy load non-critical images

## J. Enable Compression

Netlify automatically handles this, but verify in `netlify.toml`:

```
[[headers]]
for = "/*"
[headers.values]
Content-Encoding = "gzip, br"
```

## 3. Monitoring & Scaling Strategy

### Recommended Monitoring

1. **Netlify Analytics:** Track concurrent users and bandwidth
2. **Supabase Dashboard:** Monitor database connections and query performance
3. **Browser Performance API:** Track client-side performance

```
// Add to App.tsx
useEffect(() => {
  const perfData = performance.getEntriesByType('navigation')[0];
  console.log('Page load time:', perfData.loadEventEnd -
perfData.fetchStart);
}, []);
```

### Scaling Triggers

Metric	Warning Threshold	Action
Concurrent DB connections	> 400	Implement caching (Optimization A)
Page load time	> 3 seconds	Enable service worker (Optimization D)
Bandwidth usage	> 80%	Optimize assets (Optimization I)

## 4. Cost-Benefit Analysis

Optimization	Implementation Time	Impact	Priority
A. Question Caching	2-3 hours	🔥 Very High	P0
B. Prefetch Questions	1 hour	🔥 High	P0
C. Database Query Optimization	2 hours	🔥 High	P1
D. Service Worker	3-4 hours	🔥 High	P1
E. Lazy Load Audio	1-2 hours	● Medium	P2
F. Canvas Optimization	4-5 hours	● Medium	P2
G. Connection Pooling	✓ Already done	-	-
H. Code Splitting	1 hour	● Low	P3
I. Image Optimization	2 hours	● Low	P3

## 5. Summary & Recommendations

### Current State

- ✓ **Good:** Client-side architecture minimizes server load
- ✓ **Good:** Static hosting on CDN
- ✓ **Good:** Indexed database queries
- ⚠ **Concern:** No caching strategy
- ⚠ **Concern:** Database queries on every game start

### Immediate Actions (Next 1-2 Days)

1. **Implement Question Caching** (Optimization A) - **2-3 hours**
2. **Add Question Prefetching** (Optimization B) - **1 hour**

**Expected Result:** Capacity increases from **500-1,000** to **5,000-10,000** concurrent users.

### Short-term Actions (Next 1-2 Weeks)

3. **Optimize Database Queries** (Optimization C) - **2 hours**
4. **Add Service Worker** (Optimization D) - **3-4 hours**

**Expected Result:** Near-instant load times, offline support, minimal server load.

### Long-term Considerations

- Monitor actual usage patterns
- Consider upgrading Supabase tier if database becomes bottleneck
- Implement analytics to track real concurrent user counts

## Conclusion

**Your current setup can handle 500-1,000 concurrent users** with the primary bottleneck being Supabase database connections, not your Netlify server.

**With recommended optimizations (especially A & B), you can scale to 5,000-10,000+ concurrent users** without any infrastructure changes.

The beauty of your architecture is that **most processing happens client-side**, so your server resources are barely utilized. Focus on reducing unnecessary database queries through caching, and you'll be able to handle massive scale.